



Machine Learning

LABORATORY: REGRESSION

NAME: 張子中

STUDENT ID#: 313605013

Objectives:

- This assignment aims to develop a deeper understanding of linear regression by exploring their mathematical foundations, implementation, and evaluation using gradient descent.
- Understand the concepts and mathematics behind regression models.
- Implement and Evaluate regression models **from scratch**.
- Understand the key concepts of regression, including least squares, likelihood estimation, and bias–variance trade-off.

Part 1. Background

Here we will know the basics concepts that we will use for the implementation of this algorithm.

What is Regression? Regression is a statistical approach to model the relationship between a dependent variable (output) and one or more independent variables (inputs). Linear Regression is used to solve problems where the relationship between variables can be reasonably approximated by a straight line.

Tasks: In this assignment, you are required to implement linear and logistic regression models using only NumPy. You will get no points by simply calling `sklearn.linear_model.LinearRegression`. Your task is to train these models with gradient Descent on a provided dataset, evaluate their performance, and test them on unseen data. The dataset and sample code can be found here: <https://github.com/Satriosnjya/ML-Labs.git>

Part 2. Arithmetic Instructions.

Step	Procedure
1	Define the Model: Linear regression predicts the output y using a linear function : $y(x, w) = w_0 + w_1x_1 + \dots + w_Dx_D$, In simple linear regression with only one feature x : $y = mx + b$
2	Mean Square Error (MSE) Loss Calculation From the textbook, <i>Equation (4.11)</i> defines the sum-of-squares error function: $E_D(w) = \frac{1}{2} \sum_{n=1}^N (t_n - w^T \phi(x_n))^2$ which directly follows the equation, computing the squared difference between predicted and actual values.
3	Gradient Descent for Weight Updates From the textbook, <i>Equation (4.12)</i> defines the gradient of the log-likelihood: $\nabla_w \ln p(t X, w, \sigma^2) = \frac{1}{\sigma^2} \sum_{n=1}^N (t_n - w^T \phi(x_n)) \phi(x_n)^T$ This follows the principle of gradient descent , adjusting weights by subtracting a scaled version of the gradient.
4	Model Training Using Gradient Descent Equation Reference: From the textbook, <i>Equation (4.22)</i> for sequential learning (LMS Algorithm):



$$w^{(\tau+1)} = w^{(\tau)} + \eta(t_n - w^T \phi(x_n)) \phi_n$$

Follows this equation by iteratively updating w based on the gradient.

5 Mean Square Error (MSE) as Evaluation Metric

From the textbook, *Equation (4.20)* defines the MLE estimate for variance:

$$\sigma_{ML}^2 = \frac{1}{N} \sum_{n=1}^N (t_n - w_{ML}^T \phi(x_n))^2$$

Part 3. Data Transfer Instructions.

Step	Procedure
1	Download dataset and example code from https://github.com/Satriosnjva/ML-Labs.git
2	Open colab.research.google.com or you can run the python code in your own computer
3	Colab : Make a new Notebook, connect the runtime, then upload <code>regression_data.npy</code>
4	Load Libraries <i>Load Libraries and Generate Data</i> <code>import numpy as np</code> <code>import matplotlib.pyplot as plt</code>
5	Generate Data <code># Load dataset</code> <code>x_train, x_test, y_train, y_test = np.load('regression_data.npy', allow_pickle=True)</code> <code># Reshape targets</code> <code>y_train = y_train.reshape(-1,)</code> <code>y_test = y_test.reshape(-1,)</code> <code># Add bias term (column of ones)</code> <code>train_data = np.hstack((x_train, np.ones((x_train.shape[0], 1))))</code> <code>test_data = np.hstack((x_test, np.ones((x_test.shape[0], 1))))</code>

Grading & Submission Instructions

Hands-on Tasks:

- (10%) Implement Standard Linear Regression using Gradient Descent
 - Compute gradients for weight (m) and bias (b).
 - Update weights.
 - Output : Model parameters (weight and bias)
- (10%) Evaluate the model using MSE for standard regression
 - Complete `compute_mse()` function
 - Output : MSE for standard regression
- (10%) Implement Ridge Regression with L2 Regularization
 - Modify loss function to include L2
 - Compute updated gradients for weight and bias.
 - Output : Model parameters and MSE for Ridge Regression
- (10%) Plot the training loss curve.
 - Store the loss at each iteration and plot it using `matplotlib`.
 - Try different values for: Learning Rate (lr) and Number of Iterations (*iterations*)
 - Output : Loss curve comparison of Standard regression and ridge regression

Assignment:

- (20%) Implement Closed-form Ridge Regression (Refer to Equation 4.27)

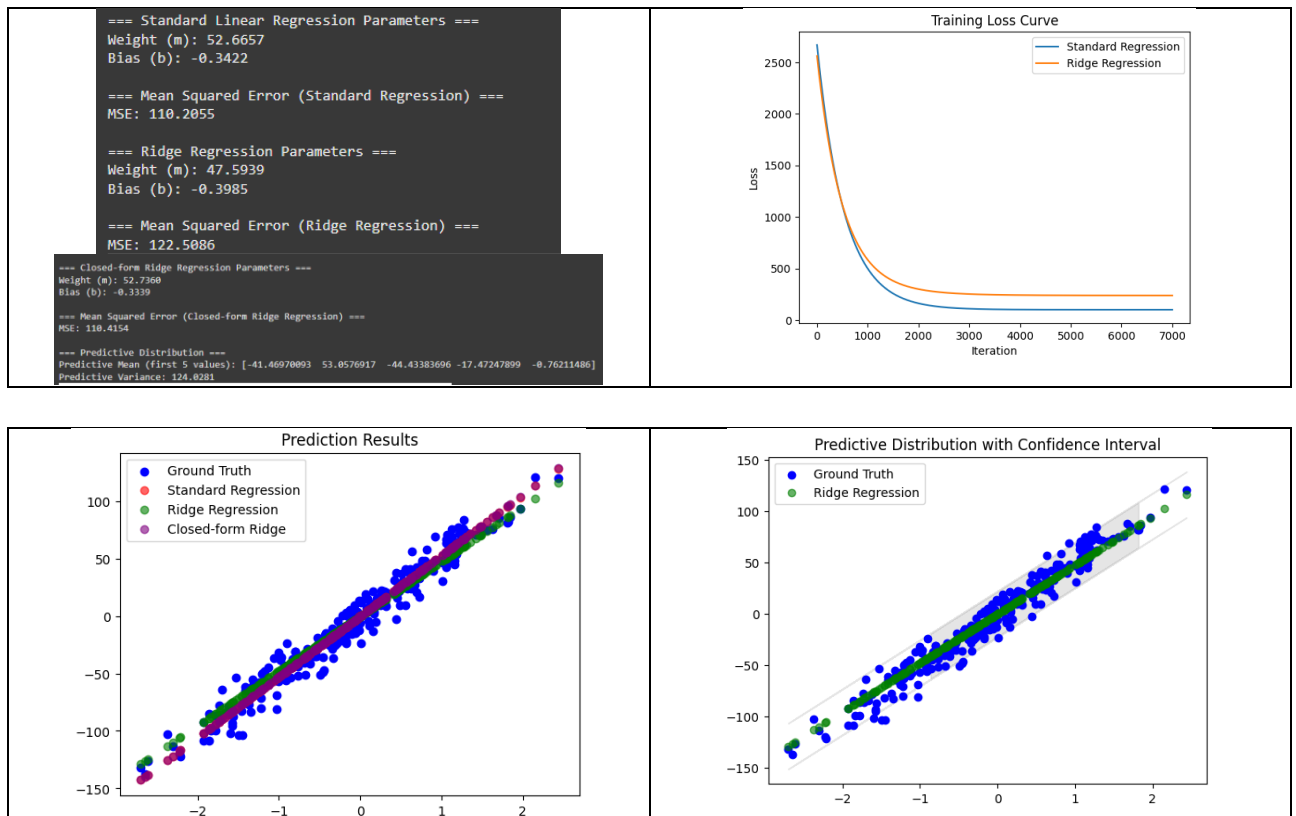


- a. Fill in `closed_form_ridge()` function
 - b. Compute `w_closed_form`
 - c. Output : Model parameters and MSE for closed-form Ridge Regression
6. (20%) Implement predictive distribution (Refer to Equation 4.33)
 - a. Fill in `predictive_mean` and `predictive_variance`
 - b. Print 5 values of `predictive_mean`
 - c. Output : Predictive variance & example predictions
7. (20%) Visualize predictions and confidence intervals
 - a. Fill in missing values in scatter plot (`plt.scatter`)
 - b. Implement confidence interval shading using `plt.fill_between()`.
 - c. Output :
 - i. Prediction plot of standard regression, ridge regression and closed-form regression.
 - ii. Predictive distribution with confidence interval.

Submission:

1. Report: Include screenshots of your results for each task (model training, evaluation, and plots) in the last pages of this PDF file.
2. Code: Submit your complete Python script (.py or .ipynb notebook).
3. Upload both your report and code to the E3 system. Name your files correctly:
 - a. Report: StudentID_Lab1.pdf
 - b. Code: StudentID_Lab1.py or StudentID_Lab1.ipynb
4. 1 day late: 10% deduction from total score. (Due Date : Sunday 9:00 PM)
5. Plagiarism is **strictly prohibited**. Submitting copied work from other students will result in penalties.

Sample Output: Please note that it is for reference only



Put Your Code Results Here:

```
=== Standard Linear Regression Parameters ===
Weight (m): 52.74342236734251
Bias (b): -0.33377936258693275
```

```
=== Mean Squared Error (Standard Regression) ===
MSE: 110.43782679931013
```

```
=== Ridge Regression Parameters ===
Weight (m): 47.627210226184786
Bias (b): -0.4344245396430874
```

```
=== Mean Squared Error (Ridge Regression) ===
MSE: 122.20206906667423
```

```
=== Closed-form Ridge Regression Parameters ===
Weight (m): 52.735986986264884
Bias (b): -0.3339074989362677
```

```
=== Mean Squared Error (Closed-form Ridge Regression) ===
MSE: 110.41535937243957
```

```
=== Predictive Distribution ===
Predictive Mean (first 5 values): [-45.84242019 58.89769953 -49.12680109 -19.25254154 -0.73679187]
Predictive Variance: 99.62103620567123
```

